

Full Name:

Andrew Id:

15-418/618 Spring 2020

Exercise 2

Assigned: Fri., Jan. 31

Due: Fri., Feb. 7, 11:00 pm

Overview

This exercise is designed to help you better understand the lecture material and be prepared for the style of questions you will get on the exams. The questions are designed to have simple answers. Any explanation you provide can be brief—at most 3 sentences. You should work on this on your own, since that's how things will be when you take an exam.

You will submit an electronic version of this assignment to Gradescope as a PDF file. For those of you familiar with the \LaTeX text formatter, you can download the template and configuration files at:

<http://www.cs.cmu.edu/~418/exercises/config-ex2.tex>

<http://www.cs.cmu.edu/~418/exercises/ex2.tex>

Instructions for how to use this template are included as comments in the file. Otherwise, you can use this PDF document as your starting point. You can either: 1) electronically modify the PDF, or 2) print it out, write your answers by hand, and scan it. In any case, we expect your solution to follow the formatting of this document.

Problem 1: Instruction-Level Parallelism

The following set of problems concern instruction-level parallelism and the limitations to performance of loop code imposed by data dependencies and resource limitations (as expressed by throughput bounds).

Calculating the distance between two vectors is an important operation for many linear algebra packages and can be defined as:

$$\sqrt{\sum_{i=0}^{N-1} (A[i] - B[i])^2}.$$

Consider the following C code for calculating the squared distance between two vectors.

```
float distance(float A[], float B[], int N) {
    int i = 0;
    float total = 0;
    while (i < N) {
        float a = A[i];           // Load
        float b = B[i];           // Load
        float diff = a - b;       // FP Add
        float squared = diff * diff; // FP Multiply
        total = total + squared;  // FP Add
        ++i;                     // Performed by integer unit
    }
    return total;
}
```

Suppose for the following questions you have a machine with multiple execution units of the following types:

Load/Store: Performs load and store operations. Can perform its own address arithmetic.

Floating-point Adder: Performs floating-point addition and subtraction

Floating-point Multiplier: Performs floating-point multiplication

Integer: Performs integer operations, including addition and comparison.

The processor has the following combination of execution units:

Unit	Count	Latency
Load/Store	1	4
Floating-point Add	2	1
Floating-point Multiply	1	2
Integer Add	2	1

Each of the multi-cycle units is *fully pipelined*, able to begin a new operation on each clock cycle.

You should also assume the following:

- The compiler produces optimized code. All local variables are held in registers.
- The machine described has a sophisticated CPU with support for out-of-order execution, pipelining, branch prediction, and speculation.
- N is very large ($> 10^6$).
- There are no cache misses.
- The only limits to average program performance are 1) the latencies due to data dependencies on loop-carried variables, and 2) the throughput limits of the functional units.

- A. Draw the dataflow diagram for the operations in the loop of `distance`. Use the local variable names as labels.

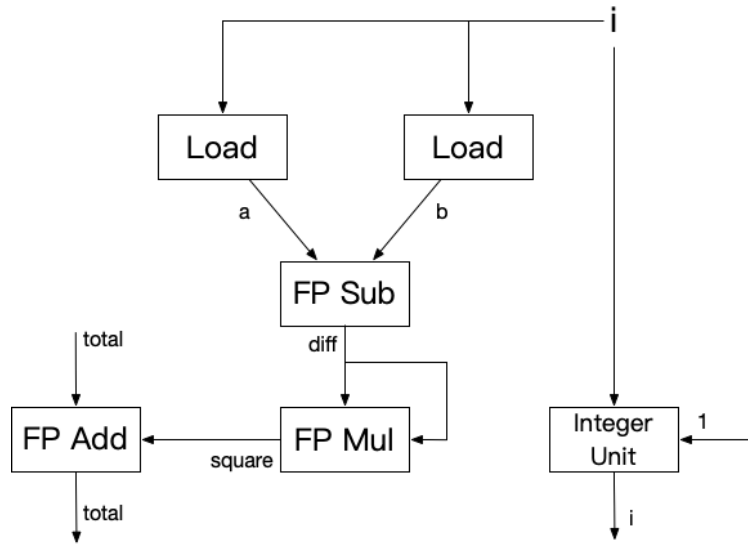


Figure 1: The dataflow diagram

- B. Which local variables create loop dependencies?

Total and `i` create loop dependencies.

- C. What latency bound does the loop of `distance` impose on the minimum average number of cycles for each element computed? Justify your answer.

A latency bound is encountered when a series of operations must be performed in strict sequence. If there is some chain of data dependencies in a program where the sum of all of the latencies along that chain equals T , then the program will require at least T cycles to execute.

Here, only the dependencies between iterations of the loop is considered. The addition of `total` and `squared`, and the increment of `i` both create latency bounds of 1. These are the only two calculations that depends on the previous iterations, and must be performed in sequence.

- D. What throughput bound does the set of available execution units impose on the loop of `distance`? Explain.

A throughput bound characterizes the raw computing capacity of the processor's functional units. Assume that a program requires a total of N computations of some operation, that the processor has C functional units capable of performing that operation, and that these units have an issue time of I . Then the program will require at least $N \cdot I/C$ cycles to execute. A program can achieve the throughput bound only when it can keep the pipelines filled for all of the functional units.

The 2 Load instructions share a Load/Store unit, and thus create 2 cycles. All other functional units (FP Add, FP Mul, integer) create a bound of 1 cycle each.

- E. If you could add one additional execution unit to the mix above, what would it be? How would that change your program performance?

I would add another Load/Store Unit. This could reduce one cycle of the throughput bound. And matches the latency bound.

Barrier Synchronization

The following set of problems are inspired by the code shown in Slides 48 and 49 in Lecture 05. In the following, we show only key parts of our version of the code. You can get the complete versions in the directory

<http://www.cs.cmu.edu/~418/exercises/ex2-code>

We suggest you download this code and study it. You can also compile and run it. (See the `README.txt` file.) Insert print statements into the code to trace the actions of the different threads.

The synchronization code in the slides show implement parts of a grid solver. To focus more directly on synchronization issues, we will adapt the code for the following, highly contrived, application.

Let B denote a *batch size* and θ , with $0.0 < \theta < 1.0$ denote a *threshold*. Suppose we perform a series of *phases*, where in each phase we compute B random numbers, each ranging between 0.0 and 1.0, and take their average a . How many phases will it take to reach a case where $a \leq \theta$, and what is the achieved value of a ? We assume the “random” numbers are actually generated by a pseudo-random generator, and so, assuming we always start with the same seed, the results will be deterministic.

Our implementation has B threads running concurrently, using barrier synchronization to keep them operating on the same phase. The global variables are as follows:

```
//// Global variables ////
// Read-only
float target_average; // Convergence goal
int batch_size;      // Number in batch
// Only written by single thread
long pcount = 0;     // Number of phases required
// Read-write
float phase_sum;      // Sum of all values in current phase
```

A Three-Barrier Solution

The following is the thread procedure for a version using three barriers, similar to the code in Slide 48. The full program is in the file `rconverge1.c`. There are barriers between the three actions performed on shared variable `phase_sum` in each phase: setting it to zero, incrementing by the random value, and testing for convergence. The call `uniform()` returns a (pseudo-)random number between 0.0 and 1.0.

```

//// Thread procedure # 1 ////
void *thread_proc(void *ival) {
    long myid = (long) ival;
    bool myconverged = false;
    long count = 0;
    while (!myconverged) {
        count++;

        phase_sum = 0.0;           // Action 1: Set to zero

        barrier();    // Barrier #1

        float myval = uniform();
        atomic_add(&phase_sum, myval); // Action 2: Increment

        barrier();    // Barrier #2

        myconverged =           // Action 3: Test
            (phase_sum/batch_size) <= target_average;

        barrier();    // Barrier #3

    }
    if (myid == 0)
        pcount = count;
    return NULL;
}

```

Problem 2: Understanding the Three-Barrier Solution

For each of the three barriers, what can go wrong if you eliminate it? You can try this with the actual code. You may want to insert print statements to track what happens. Describe the behavior you observe, and explain why it is happening.

A. Barrier #1

When barrier #1 is eliminated, different threads will set `phase_sum` to 0 at different time, which might cause the loss of random values that have already added to `phase_sum`. So the final `phase_sum` is not the sum of all random values, and thus the program will converge easier.

B. Barrier #2

When barrier #2 is eliminated, different threads judge `myconverged` with different values of `phase_sum` since they can reach the instruction at any time they want. So different threads can have different values of `myconverged`, and thus some threads might return before others. However, when some threads have not yet returned, they can never get off the barrier #1 in the next phase, because there is not as much threads as required. And the whole program will keep waiting.

C. Barrier #3

When barrier #2 is released, all threads have the same `phase_sum`. If that phase have not yet converged, `myconverged` should be `false`. But without barrier #3, some threads might begin the next phase first, and reset `phase_sum` to 0, and cause the threads that is still in the last phase get a wrong `myconverged` due to a wrong `phase_sum = 0`. And the rest is similar to the elimination of Barrier #2.

Single-Barrier Solutions

As Slide 49 demonstrates, it is possible to transform the three-barrier solution into one that uses only one barrier, using a technique known as *software pipelining*. The idea is to maintain multiple versions of the shared, global variables, and make sure that the different actions being performed within the loop operate on different versions. The number of versions required is referred to as the *pipeline depth*.

For our case, we only need to have multiple copies of the variable `phase_sum`:

```
float phase_sum[PIPEDEPTH]; // Sum of all values in current phase
```

The quantity PIPEDEPTH is a compile-time constant. In the version provided to you, it is set to 128.

First, we will explore a version of the thread procedure that is easier to understand. The thread procedure is shown below. The full code is in the file `rconverge2.c`. We see that it maintains three indices into the `phase_sum`, one for each of the basic operations. These indices are incremented by one for each phase, modulo the pipeline depth.

```
//// Thread procedure #2 ////
void *thread_proc(void *ival) {
    long myid = (long) ival;
    bool myconverged = false;
    int previndex = 0;
    int index = 1;
    long count = 0;
    phase_sum[1] = 0;
    barrier();          // Barrier #1

    while (!myconverged) {
        count++;

        int nextindex = (index+1) % PIPEDEPTH;
        phase_sum[nextindex] = 0.0;          // Action 1: Set to zero

        float myval = uniform();
        atomic_add(&phase_sum[index], myval); // Action 2: Increment

        myconverged =                          // Action 3: Test
            (phase_sum[previndex]/batch_size) <= target_average;

        barrier();    // Barrier #2

        previndex = index;
        index = nextindex;
    }
    if (myid == 0)
        pcount = count - 1;
    return NULL;
}
```


Problem 3: Understanding the First Single-Barrier Solution

- A. Barrier #1 is called only at the beginning of the procedure. Describe its purpose. What incorrect behavior can arise by eliminating it?

The idea of this version of one barrier algorithm is that we check if the last phase is converged, calculate the sum of the current phase, and set the `phase_sum` of the next phase to 0. Barrier #2 is meant to make sure all threads go to the next phase together, since the current phase will overwrite the data of next phase.

The purpose of barrier #1 is to make sure all threads go to phase #1 together. Without the barrier #1, the slower threads will overwrite the addition of those faster threads when they executing `phase_sum[1] = 0`.

- B. Why does global variable `pcount` get set to `count - 1` in this code, but to `count` in the earlier code?

Because the `myconverged` test only test whether the last phase has converged, so if it passed, it means that the convergence is achieved in the last phase, instead of the current phase.

- C. How small can constant `PIPEDEPTH` be set and still get correct behavior? Explain why this is the case.

The smallest `PIPEDEPTH` could be 3. In each phase, we will concern about three phases, the previous one, the current one, and the next one. We only need to make sure that those three phases are in different position in `phase_sum`.

- D. Suppose you swap the order of Barrier #2 with the code implementing Action #3. Explain why the program still works for the default version of `PIPEDEPTH`.

When we do so, the threads might be different phase with each other when testing and setting to 0. Some threads are testing the previous phase of phase i , while some are setting the next phase of phase $i + 1$ to zero. Namely, some threads are testing phase $i - 1$, while some are setting phase $i + 2$.

When the `PIPEDEPTH` is 128, the subscripts of phase $i - 1$ and phase $i + 2$ are different, so there is no conflict.

E. With this swapped version, how small can you set PIPEDEPTH and still get the correct answer? Explain.

We need to ensure that the subscripts of $i - 1$ and $i + 2$ is different.

When the PIPEDEPTH is 3, those two subscripts are the same, and thus conflict will be caused. If some threads testing after other threads set the `phase_sum` to 0, they will return earlier, and leave other threads waiting in the barrier #2 of next phase.

The smallest number to do so is 4.

Problem 4: Understanding the Second Single-Barrier Solution

Our final version of the code resembles that seen in Slide 49:

```
//// Thread procedure #3 ////
void *thread_proc(void *ival) {
    long myid = (long) ival;
    bool myconverged = false;
    int index = 1;
    long count = 0;
    phase_sum[1] = 0;
    barrier();          // Barrier #1

    while (!myconverged) {
        count++;

        int nextindex = (index+1) % PIPEDEPTH;
        phase_sum[nextindex] = 0.0;          // Action 1: Set to zero

        float myval = uniform();
        atomic_add(&phase_sum[index], myval); // Action 2: Increment

        barrier();    // Barrier #2

        myconverged =          // Action 3: Test
            (phase_sum[index]/batch_size) <= target_average;

        index = nextindex;
    }
    if (myid == 0)
        pcount = count;
    return NULL;
}
```

This version maintains only two indices, and the barrier synchronization has been put before Action #3. Let us explore this code:

A. Explain why the code does not require the third index `previndex`

In the last version, we test on the previous phase because barrier #2 is after the test, and we cannot ensure that all additions in this phase are finished when testing. However, in this version, all threads have passed barrier #2 when testing, and thus we know that all additions have finished. So the `phase_sum` of the current phase is the final sum wanted when testing. Thus, there is no need to check the current phase in the next phase.

B. In this version, the global variable `pcount` is set to the local value `count`, without decrementing it. Explain why this is the correct result.

As aforementioned, the testing is for the current phase, not the previous phase. So when `myconverged = true`, that means the current phase is the answer.

C. How small can constant `PIPEDEPTH` be set and still get correct behavior? Explain why this is the case.

This is a little similar to 3D and 3E. There might exist a situation that the slower threads are testing the current phase of phase i while others are setting the next phase of phase $i + 1$ to zero. Namely, some threads are testing the phase i while others are setting phase $i + 2$ to 0.

We need to ensure that the subscripts of phase i and phase $i + 2$ have no conflict. So the smallest `PIPEDEPTH` is 3.

Cilk Scheduling

The following problems are based on the presentation of the Cilk programming environment from Lecture 06.

Suppose we are running a program that uses the Cilk mechanisms for fork-join parallelism. Assume the following:

- The system is running two threads
- Every execution of `cilk_spawn` requires 1 millisecond. The executing thread will push its continuation at the top of its queue and begin executing (after 1ms) the spawned function (“child first” scheduling).
- One thread can steal work from the queue of another. It always steals from the bottom of the queue. Stealing requires 2 milliseconds, and then the thread can begin performing whatever operation is specified in the record.
- The procedure `foo` requires 3 milliseconds to complete.
- All other operations require only a negligible amount of time.

As an example, consider the following code

```
void simple() {  
    cilk_spawn foo(1);  
    foo(3);  
    cilk_sync;  
}
```

If Thread 1 starts executing `simple`, we can trace its execution as follows:

Time	Thread 1	Thread 2
0	Spawn <code>foo(1)</code> ; Push <code>foo(3)</code>	Idle
1	Execute <code>foo(1)</code>	Steal <code>foo(3)</code>
2	Executing	Stealing
3	Executing	Execute <code>foo(3)</code>
4	Idle	Executing
5	Idle	Executing

Problem 5: Divide and Conquer Parallelism

Consider the following function for executing multiple copies of function `foo` via a series of forks that repeatedly split the problem in half, analogous to the control structure of quicksort:

```
void rfor(int start, int last) {
    if (start == last)
        foo(start);
    else {
        int middle = (start + last)/2;
        cilk_spawn rfor(start, middle);
        rfor(middle+1, last);
    }
    cilk_sync;
}
```

- A. Fill in the following table, showing how two threads would handle the execution of `rfor(0, 3)`. (You may not require all of the rows in the table.)

Time	Thread 1	Thread 2
0	Spawn <code>rfor(0, 1)</code> , push <code>rfor(2, 3)</code>	Idle
1	Spawn <code>rfor(0, 0)</code> ; push <code>rfor(1, 1)</code>	Steal <code>rfor(2, 3)</code>
2	Executing <code>foo(0)</code>	Stealing
3	Executing	Spawn <code>rfor(2, 2)</code> ; push <code>rfor(3, 3)</code>
4	Executing	Executing <code>foo(2)</code>
5	Executing <code>foo(1)</code>	Executing
6	Executing	Executing
7	Executing	Executing <code>foo(3)</code>
8	Idle	Executing
9	Idle	Executing

B. What do you get as the computed speedup for this execution?

Serial execution takes $4 \times 3 = 12$ ms, while parallel execution takes 10 ms as shown above.
The speedup = $12/10 = 1.2$

Problem 6: Iterative Parallelism

Consider the following function for executing multiple copies of function `foo` by having the main thread spawning off series of threads, each executing one instance of `foo`. (Although the function is written as a recursive procedure, you will see that the sequence of spawns is identical what would occur if they were done as a single loop, as in shown in Lecture 06, starting with Slides 39.)

```
void ifor(int start, int last) {
    if (start == last)
        foo(start);
    else {
        cilk_spawn ifor(start, start);
        ifor(start+1, last);
    }
    cilk_synch;
}
```

A. Fill in the following table, showing how two threads would handle the execution of `ifor(0, 3)`:

Time	Thread 1	Thread 2
0	Spawn <code>ifor(0, 0)</code> ; push <code>ifor(1, 3)</code>	Idle
1	Execute <code>foo(0)</code>	Steal <code>ifor(1, 3)</code>
2	Executing	Stealing
3	Executing	Spawn <code>ifor(1, 1)</code> ; push <code>ifor(2, 3)</code>
4	Steal <code>ifor(2, 3)</code>	Execute <code>foo(1)</code>
5	Stealing	Executing
6	Spawn <code>ifor(2, 2)</code> ; push <code>ifor(3, 3)</code>	Executing
7	Execute <code>foo(2)</code>	Steal <code>ifor(3, 3)</code>
8	Executing	Stealing
9	Executing	Executing <code>foo(3)</code>
10	Idle	Executing
11	Idle	Executing

B. What do you get as the computed speedup for this execution?

The parallel execution takes 12 ms. So the speedup = $12/12 = 1.0$

C. What insights do these to example give you regarding the best way to use Cilk in exploiting fork-join parallelism?

Stealing is costly, so we want to reduce stealing as much as possible. Compare to decrease-and-conquer in problem 6, divide-and-conquer in problem 5 is better.